

Geavanceerde Computerarchitectuur

3rd session

Optimization techniques,
Partitioning the grid into warps and reducing thread divergence

Thomas Feys, thomas.feys@kueluven.be

DRAMCO – WaveCore – ESAT

Academic year 2024/25

0 Outline

① Background

② Exercise

1 Outline

① Background

② Exercise

1 A warp

- ▶ A warp is a group of 32 threads
- ▶ number of warps = $\left\lceil \frac{\text{Threads per block}}{32} \right\rceil \cdot \text{number of blocks}$
- ▶ For example
 - `foo<<<2,64>>>()`
 - `foo<<<4,33>>>()`

1 A warp

- ▶ A warp is a group of 32 threads
- ▶ number of warps = $\left\lceil \frac{\text{Threads per block}}{32} \right\rceil \cdot \text{number of blocks}$
- ▶ For example
 - `foo<<<2,64>>>()` \Rightarrow 4 (128 threads ... 128 cores)
 - `foo<<<4,33>>>()`

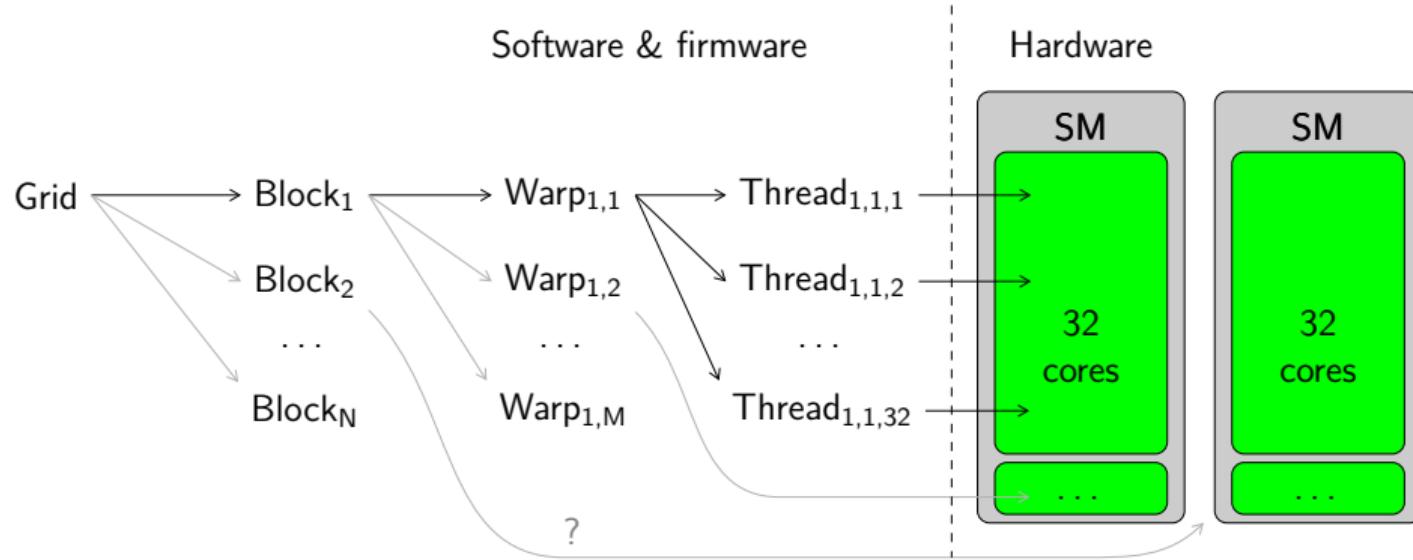
1 A warp

- ▶ A warp is a group of 32 threads
- ▶ number of warps = $\left\lceil \frac{\text{Threads per block}}{32} \right\rceil \cdot \text{number of blocks}$
- ▶ For example
 - `foo<<<2,64>>>()` \Rightarrow 4 (128 threads ... 128 cores)
 - `foo<<<4,33>>>()` \Rightarrow 8 (132 threads ... **256 cores**)

1 A warp

- ▶ A warp is a group of 32 threads
- ▶ number of warps = $\left\lceil \frac{\text{Threads per block}}{32} \right\rceil \cdot \text{number of blocks}$
- ▶ For example
 - `foo<<<2,64>>>()` \Rightarrow 4 (128 threads ... 128 cores)
 - `foo<<<4,33>>>()` \Rightarrow 8 (132 threads ... **256 cores**)
- ▶ Also called a *sub-group* in OpenCL, a *wavefront* by AMD, ...

1 Threads are grouped into warps and mapped to cores



A block is divided into warps. The warps of a block are associated with a single stream multiprocessor (SM) – the GT 740 has 2 SMs, each with 192 cores.

1 "Warping" can over provision GPU resources

Assume an input array with 384 elements and a GPU with 384 CUDA cores
How many warps are spawned?

- ▶ `foo<<<2, 192>>>(...)`
- ▶ `foo<<<384, 1>>>(...)`

1 "Warping" can over provision GPU resources

Assume an input array with 384 elements and a GPU with 384 CUDA cores
How many warps are spawned?

- ▶ `foo<<<2, 192>>>(...)` \Rightarrow 12 (requires 384 cores)
- ▶ `foo<<<384, 1>>>(...)`

1 "Warping" can over provision GPU resources

Assume an input array with 384 elements and a GPU with 384 CUDA cores
How many warps are spawned?

- ▶ `foo<<<2, 192>>>(...)` ⇒ 12 (requires 384 cores)
- ▶ `foo<<<384, 1>>>(...)` ⇒ 384 (requires 12288 cores)

1 "Warping" can over provision GPU resources

Assume an input array with 384 elements and a GPU with 384 CUDA cores
How many warps are spawned?

- ▶ `foo<<<2, 192>>>(...)` ⇒ 12 (requires 384 cores)
- ▶ `foo<<<384, 1>>>(...)` ⇒ 384 (requires 12288 cores)
- ▶ `foo<<<1, 192>>>(...)` ⇒ 6 (requires 192 cores, 2 thread iterations (striding))

1 "Warping" can over provision GPU resources

Assume an input array with 384 elements and a GPU with 384 CUDA cores
How many warps are spawned?

- ▶ `foo<<<2, 192>>>(...)` ⇒ 12 (requires 384 cores)
- ▶ `foo<<<384, 1>>>(...)` ⇒ 384 (requires 12288 cores)
- ▶ `foo<<<1, 192>>>(...)` ⇒ 6 (requires 192 cores, 2 thread iterations (striding))
- ▶ `foo<<<1, 384>>>(...)` ⇒ 12 (requires 192 cores, 2 scheduling iterations)

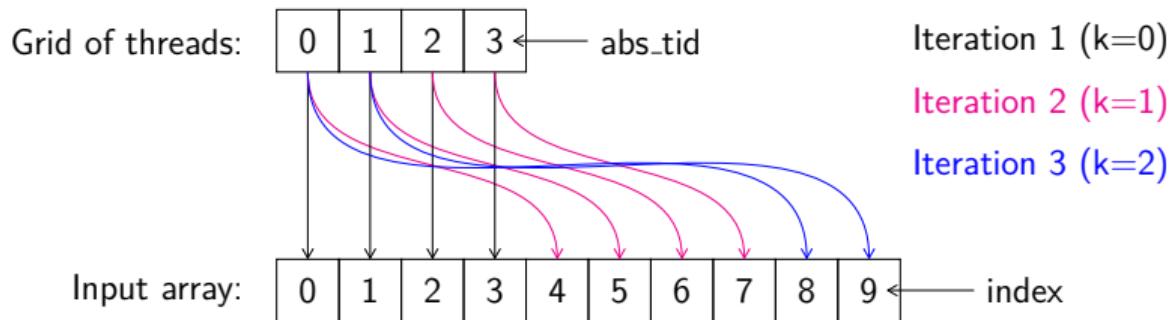
1 Thread divergence within a warp

- ▶ Part of the warp has higher complexity, taking longer to execute.
- ▶ The other part of the warp becomes idle (has finished its job).
- ▶ Idle threads occupy cores and do not contribute to the processing of data.

```
1 __global__ void addOrMultiply(int *a, int *b, int *c) {  
2     int tid = blockIdx.x * blockDim.x + threadIdx.x;  
3  
4     if (tid % 2 == 0) {          // If tid is even, add elements  
5         c[tid] = a[tid] + b[tid];  
6     } else {                  // If tid is odd, multiply elements  
7         c[tid] = a[tid] * b[tid];  
8     }  
9 }
```

1 Striding

- ▶ When the matching between threads and elements is not 1:1
- ▶ A thread may compute multiple elements
- ▶ `if(abs_tid < input_size ... ⇒ ..while(...) or ..for(...).`



```
index = abs_tid + k * grid_size , where index < input_size
```

Good explanation see:

<https://alexminnaar.com/2019/08/02/grid-stride-loops.html>

1 Striding, divergence, and warps on an example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented

1 Striding, divergence, and warps on an example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented
- ▶ Is there thread divergence?

1 Striding, divergence, and warps on an example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented

- ▶ Is there thread divergence? Yes, at least one warp of 32 threads has to run twice, while only 4 threads are needed. Divergence can be avoided by using inputs with a size that is a multiple of 32.

1 Striding, divergence, and warps on an example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented

- ▶ Is there thread divergence? Yes, at least one warp of 32 threads has to run twice, while only 4 threads are needed. Divergence can be avoided by using inputs with a size that is a multiple of 32.
- ▶ How many cores are utilized?

1 Striding, divergence, and warps on an example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented

- ▶ Is there thread divergence? Yes, at least one warp of 32 threads has to run twice, while only 4 threads are needed. Divergence can be avoided by using inputs with a size that is a multiple of 32.
- ▶ How many cores are utilized? Only 192 cores are utilized since a block (all of its warps) gets mapped to one SM.

1 Blackboard example

Assume we're using the GT 740 (2 SMs, each with 192 cores) and:

- ▶ We spawn a kernel with a single block of 256 threads
- ▶ To process an input array with 260 elements
- ▶ Striding is implemented

2 Outline

① Background

② Exercise

2 Goal

Part 1 – warps

- ▶ Implement a kernel that inverts an image (calculate $x = 255 - x$, where $x \in (0, 255)$, for each color component – RGB).
- ▶ Assess how the number of threads (full/partial warps) influences the computation time.
- ▶ See pointers below!

Part 2 – thread divergence

- ▶ Augment the inversion kernel, so that the R color component is calculated as 1D-Gaussian blur of its neighbors. Namely the i^{th} red pixel r_i is obtained as $r_i = 0.1r_{i-2} + 0.25r_{i-1} + 0.5r_i + 0.25r_{i+1} + 0.1r_{i+2}$ i.e., a Gaussian weighted sum of its 1D-neighboring red pixel values.
- ▶ Note that the green and blue values are still processed as before ($x = 255 - x$)
- ▶ Assess how thread divergence impacts the computation time.
- ▶ See pointers below!

I encourage you to discuss all parts in the report, even if not fully implemented!

2 Some pointers - nonobligatory - part 1

- ▶ Implement the inverse kernel and use the provided image load/save function (Toledo → Docs) to invert an image. You can use as many threads as there are input elements.
- ▶ Implement striding in the above kernel → if difficult, first implement striding in a simpler kernel that handles a short 1D sequence, for example, using 12 threads, 16 elements (0, 1, ..., 15), and adding 1 to each element.
- ▶ Check that the inversion kernel, with the added striding, generates a fully-inverted image even when using less threads than there are input data elements. For example, consider a ratio of 1:16 (threads:elements).

2 Some pointers - nonobligatory - part 1

- ▶ Invoke the inverse kernel from a loop, where you progressively use more threads. Time the kernel on every iteration (see cheat sheet).
- ▶ Get the average execution time, as the time may vary. For example, due to preemption by the OS.
- ▶ Note that on some setups, the GPU may take more time to finish when first invoked – you can avoid this by calling the function once before actually timing it.

2 Some pointers - nonobligatory - part 2

- ▶ Make a second kernel, where the R color component is processed according to the provided equation.
- ▶ Check that the returned image is indeed different.
- ▶ Now make a second kernel, which instead of processing `RGBRGB...RGB` operates on an input format of `RR...RGGG...GBB...B` (re-format the data correspondingly).
- ▶ Evaluate how the execution time changes with / without thread divergence.