

Geavanceerde Computerarchitectuur

4th Session

Optimization techniques,
Coalesced memory access and device (GPU) memory types

Thomas Feys, thomas.feys@kueluven.be

DRAMCO – WaveCore – ESAT

Academic year 2024/25

0 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

1 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

1 Free allocated memory

```
1 #define SIZE 16
2 ...
3 int* p1 = (int*)malloc(SIZE*sizeof(int));
4 int* p2;
5 cudaMalloc(&p2, SIZE*sizeof(int));
6 ...
7 free(p1);
8 cudaFree(p2);
```

Otherwise memory leaks are introduced.

Can use [Valgrind](#) to detect leaks.

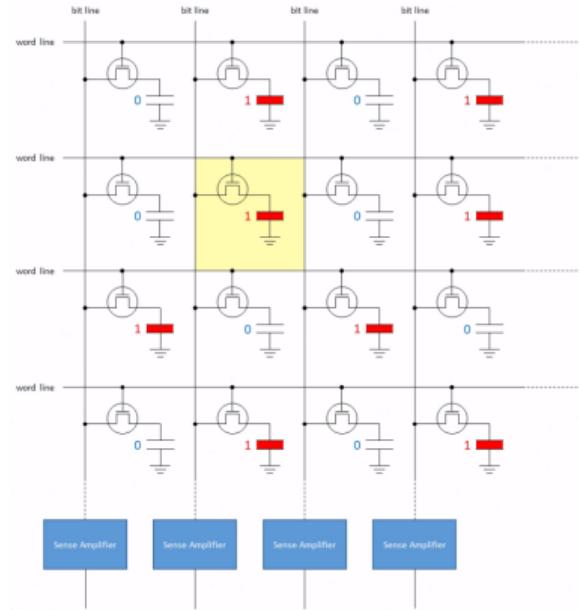
2 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

2 Coalesced memory access

Coalesce: to come together to form one larger group, substance, etc.

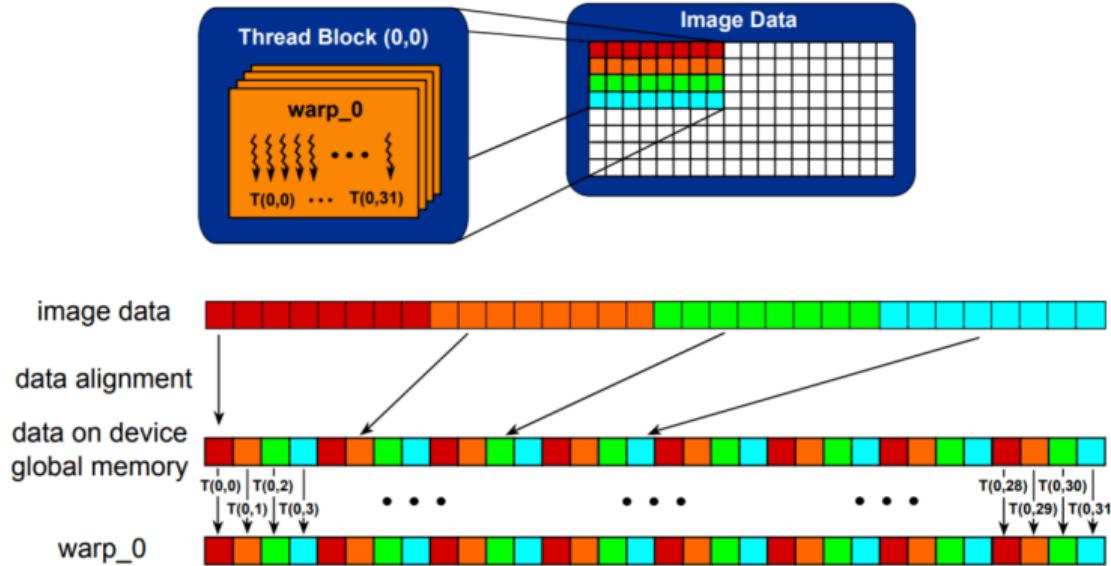
- ▶ DRAM access is done in groups (high-throughput parallel access)
- ▶ DRAM access is otherwise sequential (e.g. per request)
- ▶ Threads in a warp (32) execute the same instructions synchronously
- ▶ Group memory access reduces memory latency



2 Coalesced memory access

- ▶ Is achieved when all threads in a warp access consecutive global memory locations.
- ▶ When all threads in a warp execute a load instruction, the hardware detects whether they access consecutive global memory locations. If that's the case, the hardware combines (coalesces) all these accesses into a consolidated access to consecutive DRAM locations.
- ▶ For example, for a given load instruction of a warp, if thread 0 accesses global memory location N2, thread 1 location N+1, thread 2 location N+2, and so on, all these accesses will be coalesced into a single request for consecutive locations when accessing the DRAMs.
- ▶ Such coalesced access allows the DRAMs to deliver data as a burst.

2 Coalesced memory access in practice



2 Non-divergent threads and coalesced memory access

Non-divergent threads

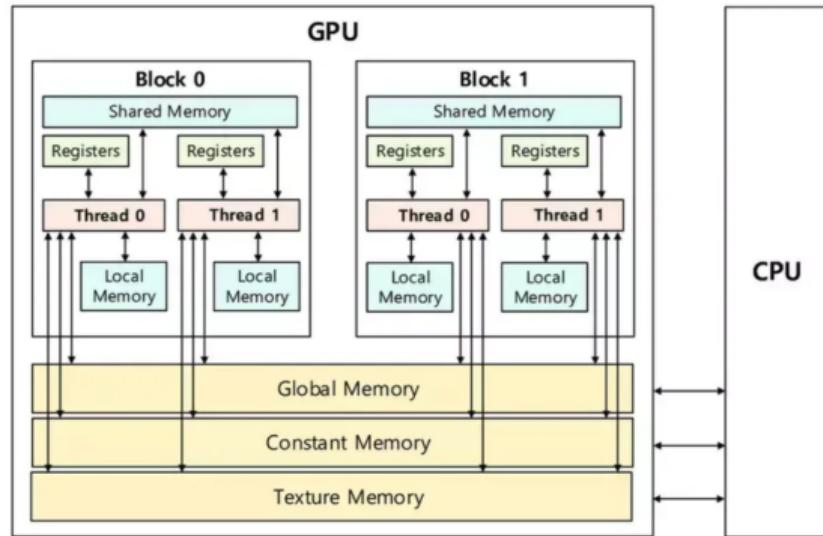
Threads in a **warp** take the same **computation** path (same complexity)

Coalesced memory access

Threads in a **warp** access the same **memory** group (adjacent addresses)

2 GPU memory types

- ▶ Shared – Per block, on-chip
- ▶ Registers – Per thread, on-chip
- ▶ Local – Per thread, off-chip
- ▶ Constant – For all, off-chip
- ▶ Global – For all, off-chip
- ▶ Texture – For all, off-chip



2 Global memory

- ▶ Off-device – high access latency
- ▶ All threads in the grid have access, as well as the CPU
- ▶ Large size – typically several GB
- ▶ The 'default' option

As in previous labs

```
1 int* garr;
2 cudaMalloc(&garr, SIZE*sizeof(int));
3 // or (see analogy to constant memory in a later slide)
4 __device__ int garr[SIZE];
```

2 Shared memory

- ▶ On-device – low access latency
- ▶ Shared among the threads of a single block
- ▶ User-managed cache, by default 48 kB (`cudaFuncSetCacheConfig(...)`)
- ▶ Two options in CUDA, see below

Hardcode `SIZE` (or use `#define`)

```
my_kernel<<<blk, tpb>>>(...); → __shared__ int smem[SIZE];
```

Define size dynamically during kernel invocation

```
my_kernel<<<blk, tpb, size>>>(...); → extern __shared__ int smem[];
```

2 Workflow with shared memory

Threads migrate data from global to shared memory

- 1 (on CPU) Copy data **to global memory** and run the kernel
- 2 (on GPU, per thread) Migrate one or more elements **to shared memory**
- 3 **Process** the data in **shared memory**
- 4 Migrate the data **back to global memory**
- 5 (on CPU) Copy data from global memory

2 Workflow with shared memory

Threads migrate data from global to shared memory

- 1 (on CPU) Copy data **to global memory** and run the kernel
- 2 (on GPU, per thread) Migrate one or more elements **to shared memory**
- 3 **Process** the data in **shared memory**
- 4 Migrate the data **back to global memory**
- 5 (on CPU) Copy data from global memory

Where should we put `--syncthreads()` ?

2 Constant memory

- ▶ Off-device, optimized for low-latency read
- ▶ All threads in the grid have access, as well as the CPU
- ▶ Small size – 64 KB

From the documentation:

```
1 __constant__ int cmem[SIZE]; // GPU
2 ...
3 int* arr; // CPU
4 ...
5 cudaMemcpyToSymbol(cmem, arr, SIZE * sizeof(int)); // From CPU to GPU
```

2 Other memory

Registers

- ▶ On-device – Low access latency
- ▶ Access per thread
- ▶ `threadIdx.x`

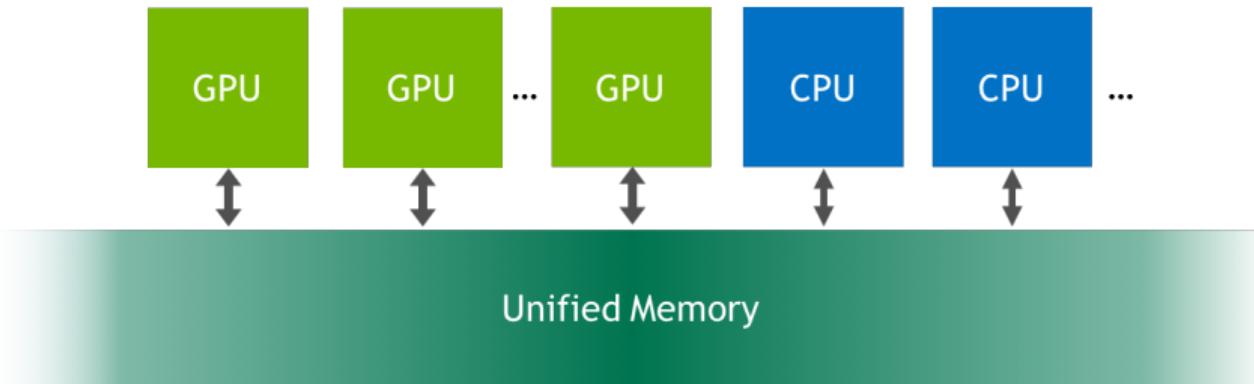
Local memory

- ▶ Off-device – High(er) access latency
- ▶ Access per thread
- ▶ `int arr[128];`

Texture memory

- ▶ Off-device with optimized read access
- ▶ Access for all threads
- ▶ For 2D graphic data

2 Regarding a question



Unified memory allows one to allocate once and use on both devices.
(Abstraction that handles memory migration.)

[More in the docs](#) and a blog post.

3 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

3 Assignment

Part 1 – Coalesced memory access

- ▶ Design a kernel that inverts the red color component in an image.
- ▶ Evaluate execution time for coalesced and un-coalesced memory access.

Part 2 – Global, shared, and constant memory

- ▶ Design a kernel that multiplies two matrices.
- ▶ Assess the processing time for a kernel that utilizes global-only, global and shared, and global and constant memory.

Evaluations should consider changing the grid and/or input size.

3 Pointers – Part 1

- ▶ Load up an image using last time's code or make a 1D array of your liking.
- ▶ Use one array which sorts data as `RGBRGB...RGB` and one as `RR...RGG...GBB...B`.
- ▶ Design two kernels, each for processing one of the above arrays.
- ▶ Calculate the inverse of the red color component, $R = 255 - R$.
- ▶ Time the kernels for different input and/or grid sizes.

3 Guiding questions – Part 1

- ▶ What is the performance difference between the coalesced and uncoalesced kernel?
- ▶ How does the performance difference change dependent on input and/or grid size?

3 Pointers – Part 2

Recall matrix multiplication:

$$2 \times 3 + 0 \times 4 = 6$$

The diagram illustrates the calculation of the element at index (0,0) of the resulting matrix from multiplying two 2x3 matrices. The first matrix has elements 2, 0, 1, 9. The second matrix has elements 3, 4, 9, 7. The result is 6. Purple ovals highlight the 2 and 3, which are multiplied together. A purple bracket above the first matrix shows the calculation $2 \times 3 + 0 \times 4 = 6$.

$$\begin{bmatrix} 2 & 0 \\ 1 & 9 \end{bmatrix} \times \begin{bmatrix} 3 & 9 \\ 4 & 7 \end{bmatrix} = \begin{bmatrix} 6 & 18 \\ 39 & 72 \end{bmatrix}$$

Image source: <https://www.codingem.com/numpy-at-operator/>

3 Pointers – Part 2

- ▶ Define an array that represents a matrix - in 1D or 2D form (start off with small square matrices and simple numbers for easier debugging).
- ▶ In this way define/allocate memory for three matrices, which will serve for $\mathbf{C} = \mathbf{A}\mathbf{B}$.
- ▶ Make a kernel that uses only global memory to compute $\mathbf{C} \rightarrow \text{verify the results.}$
- ▶ Make a second and third matrix multiplication kernel, which use shared and constant memory, respectively, in addition to global memory.
- ▶ Time the kernels for different input and/or grid sizes.

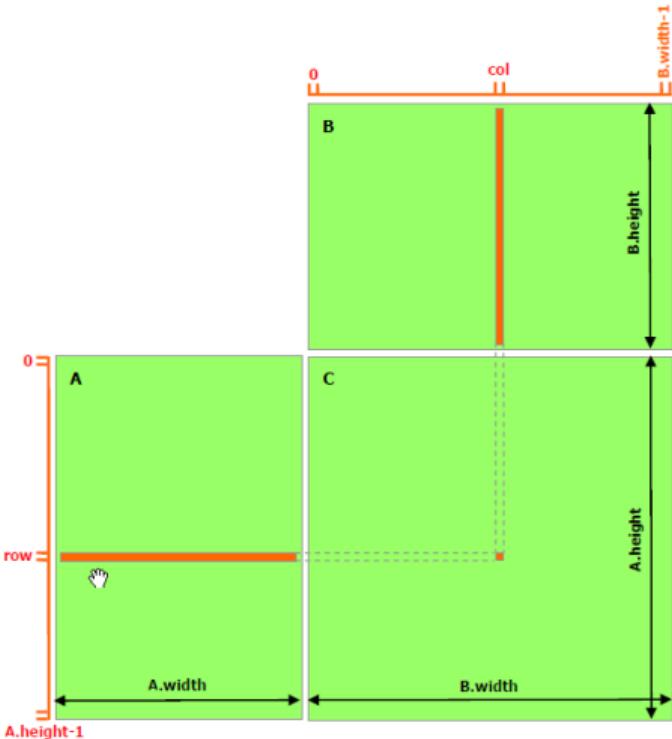
3 Guiding questions – Part 2

- ▶ How much time does memory migration take (CPU to GPU global; GPU global to GPU shared)?
- ▶ How do the execution times of matrix multiplication compare when employing the three different memory types?
- ▶ How do the times scale with input and/or grid size?

4 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

4 Memory access during matrix multiplication



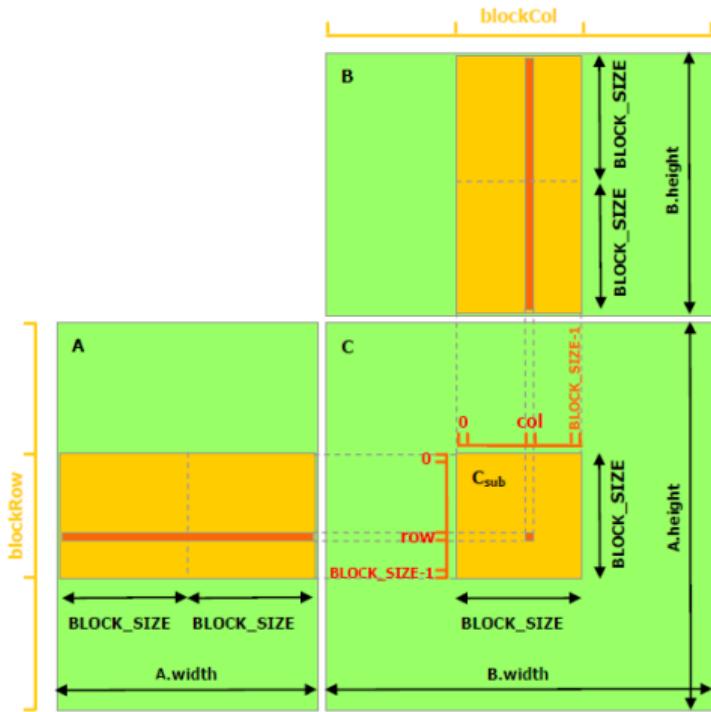
Assume:

- ▶ Matrices A, B, and C all NxN
- ▶ NxN grid of threads

Then:

- ▶ Each thread writes 1 element in C
- ▶ A-rows and B-columns are re-read N-1 times

4 Improvement by leveraging shared memory



Assume:

- ▶ Matrices **A**, **B**, and **C** all $N \times N$
- ▶ $N \times N$ grid of threads
- ▶ Square blocks with size **BLOCK_SIZE**

Then:

- ▶ Each thread writes 1 element in **C**
- ▶ Each thread copies one element from both **A** and **B** to shared memory, k -times
- ▶ **A**-rows and **B**-columns are re-read k -times
- ▶ $k = \frac{N}{BLOCK_SIZE}$

5 Outline

- ① Intermezzo
- ② GPU memory access and memory types
- ③ Exercise
- ④ Appendix
- ⑤ Sources

5 Sources

- ▶ Coalesced meaning
- ▶ DRAM intro
- ▶ Coalesced access figure
- ▶ Blog post about unified access
- ▶ Blog post about shared memory
- ▶ GPU memory figure